
Unsupervised Spectral Learning of FSTs

Raphaël Bailly Xavier Carreras Ariadna Quattoni
Universitat Politècnica de Catalunya
Barcelona, 08034
rbailly, carreras, aquattoni@lsi.upc.edu

Abstract

Finite-State Transducers (FST) are a standard tool for modeling paired input-output sequences and are used in numerous applications, ranging from computational biology to natural language processing. Recently Balle et al. [4] presented a spectral algorithm for learning FST from samples of aligned input-output sequences. In this paper we address the more realistic, yet challenging setting where the alignments are unknown to the learning algorithm. We frame FST learning as finding a low rank Hankel matrix satisfying constraints derived from observable statistics. Under this formulation, we provide identifiability results for FST distributions. Then, following previous work on rank minimization, we propose a regularized convex relaxation of this objective which is based on minimizing a nuclear norm penalty subject to linear constraints and can be solved efficiently.

1 Introduction

This paper addresses the problem of learning probability distributions over pairs of input-output sequences, also known as transduction problem. A pair of sequences is made of an *input* sequence, built from an input alphabet, and an *output* sequence, built from an output alphabet. Finite State Transducers (FST) are one of the main probabilistic tools used to model such distributions and have been used in numerous applications ranging from computational biology to natural language processing. A variety of algorithms for learning FST have been proposed in the literature, most of them are based on EM optimizations [9, 11] or grammatical inference techniques [8, 6].

In essence, an FST can be regarded as an HMM that generates bi-symbols of combined input-output symbols. The input and output symbols may be generated jointly or independently conditioned on the previous observations. A particular generation pattern constitutes what we call an alignment.

GAATTCAG-	GAATTCAG-	GAATTC-AG	GAATTC-AG
GGA-TC-GA	GGAT-C-GA	GGA-TCGA-	GGAT-CGA-

To be able to handle different alignments, a special empty symbol $*$ is added to the input and output alphabets. With this enlarged set of bi-symbols, the model is able to generate an input symbol (resp. an output symbol) without an output symbol (resp. input symbol). These special bi-symbols will be represented by the pair $\begin{smallmatrix} x \\ * \end{smallmatrix}$ (resp. $\begin{smallmatrix} * \\ y \end{smallmatrix}$). As an example, the first alignment above will correspond to the two possible representations $\begin{smallmatrix} GA*ATTTCAG* \\ G*GA*TC*GA \end{smallmatrix}$ and $\begin{smallmatrix} G*AAATTCAG* \\ GG*A*TC*GA \end{smallmatrix}$. Under this model the probability of observing a pair of un-aligned input-output sequences is obtained by integrating over all possible alignments.

Recently, following a recent trend of work on spectral learning algorithms for finite state machines [14, 2, 17, 18, 7, 16, 10, 5], Balle et al. [4] presented an algorithm for learning FST where the input to the algorithm are samples of *aligned* input-output sequences. As with most spectral methods the core idea of this algorithm is to exploit low-rank decompositions of some *Hankel matrix* representing

the distribution of aligned sequences. To estimate this Hankel matrix it is assumed that the algorithm can sample aligned sequences, i.e. it can directly observe sequences of enlarged bi-symbols.

While the problem of learning FST from fully aligned sequences (what we sometimes refer to as supervised learning) has been solved, the problem of deriving an unsupervised spectral method that can be trained from samples of input-output sequences alone (i.e. where the alignment is hidden) remains open. This setting is significantly more difficult due to the fact that we must deal with two sets of hidden variables: the states and the alignments. In this paper we address this unsupervised setting and present a spectral algorithm that can approximate the distribution of paired sequences generated by an FST without having access to aligned sequences. To the best of our knowledge this is the first spectral algorithm for this problem.

The main challenge in the unsupervised setting is that since the alignment information is not available, the Hankel matrices (as in [4]) can no longer be directly estimated from observable statistics. However, a key observation is that we can nevertheless compute observable statistics that can constraint the coefficients of the Hankel matrix. This is because the probability of observing a pair of un-aligned input-output sequences (i.e. an observable statistic) is computed by summing over all possible alignments; i.e. by summing entries of the hidden Hankel matrix. The main idea of our algorithm is to exploit these constraints and find a Hankel matrix (from which we can directly recover the model) which both agrees on the observed statistics and has a low-rank matrix factorization.

In brief, our main contribution is to show that an FST can be approximated by solving an optimization which is based on finding a low-rank matrix satisfying a set of constraints derived from observable statistics and Hankel structure. We provide sample complexity bounds and some identifiability results for this optimization that show that, theoretically, the rank and the parameters of an FST distribution can be identified. Following previous work on rank minimization, we propose a regularized convex relaxation of the proposed objective which is based on minimizing a nuclear norm penalty subject to linear constraints. The proposed relaxation balances a trade-off between model complexity (measured by the nuclear norm penalty) and fitting the observed statistics. Synthetic experiments show that the performance of our unsupervised algorithm efficiently approximates that of a supervised method trained from fully aligned sequences.

The paper is organized as follows. Section 2 gives preliminaries on FST and spectral learning methods, and establishes that an FST can be induced from a Hankel matrix of observable aligned statistics. Section 3 presents a generalized form of Hankel matrices for FST that allows to express observation constraints efficiently. One can not observe generalized Hankel matrices without assuming access to aligned samples. To solve this problem, Section 4 formulates finding the Hankel matrix of an FST from unaligned samples as rank minimization problem. This section also presents the main theoretical results of the method, as well as a convex relaxation of the rank minimization problem. Section 5 presents results on synthetic data and Section 6 concludes.

2 Preliminaries

2.1 Finite-State Transducers

Definition 1. A Finite-State Transducer (FST) of rank d is given by:

- alphabets $\Sigma_+ = \{x_1, \dots, x_p\}$ (input), $\Sigma_- = \{y_1, \dots, y_q\}$ (output)
- $\alpha_1 \in \mathbb{R}^d$, $\alpha_\infty \in \mathbb{R}^d$
- $\forall x \in \Sigma_+ \cup \{*\}, \forall y \in \Sigma_- \cup \{*\}$, a matrix $M_{xy} \in \mathbb{R}^{d \times d}$, with $M_{*} = 0$

Definition 2. Let s be an input sequence, and let t be an output sequence. An alignment of (s, t) is given by a sequence of pairs $x_1^{x_1} \dots x_n^{x_n}$ such that the sequence obtained from $x_1 \dots x_n$ (resp. $y_1 \dots y_n$) by removing the empty symbols $*$ equals s (resp. t).

Definition 3. The set of alignments for a pair of sequences (s, t) is denoted $[s, t]$.

Definition 4. Let $\Sigma = \{\Sigma_+ \cup \{*\}\} \times \{\Sigma_- \cup \{*\}\}$. The set of aligned sequences is Σ^* . The empty string is denoted ε .

Definition 5. Let T be an FST, and let $w = \frac{x_1}{y_1} \dots \frac{x_n}{y_n}$ be an aligned sequence. Then the value of w for the model T is given by:

$$r_T(w) = \alpha_1^\top M_{\frac{x_1}{y_1}} \dots M_{\frac{x_n}{y_n}} \cdot \alpha_\infty$$

Definition 6. Let (s, t) be an i/o (input/output) sequence. Then the value for (s, t) computed by an FST T is given by the sum of the values for all alignments:

$$r_T((s, t)) = \sum_{\frac{x_1}{y_1} \dots \frac{x_n}{y_n} \in [s, t]} r_T(\frac{x_1}{y_1} \dots \frac{x_n}{y_n})$$

A more complete description of FST can be found in [15].

2.2 Computing with an FST

In order to compute the value of a pair of sequences (s, t) , one needs to sum over all possible alignments, which is generally exponential in the length of s and t . Standard techniques (e.g. the edit distance algorithm) can be applied in order to compute such a value in polynomial time.

Proposition 1. Let T be an FST, $s_{1:n} \in \Sigma_+^*$, $t_{1:m} \in \Sigma_-^*$. The forward vector is defined by:


$$F_{0,j} = \alpha_1^\top M_{\frac{*}{t_1}} \dots M_{\frac{*}{t_j}}, F_{i,0} = \alpha_1^\top M_{\frac{x_i}{*}} \dots M_{\frac{x_i}{*}}, F_{i,j} = F_{i-1,j} M_{\frac{x_i}{*}} + F_{i,j-1} M_{\frac{*}{t_j}} + F_{i-1,j-1} M_{\frac{x_i}{t_j}}$$

It is then possible to compute $r_T((s, t)) = F_{n,m} \alpha_\infty$ in $O(d^2|s||t|)$. The sum of r_T over all possible values $r_T(\Sigma^*) = \sum_{s \in \Sigma_+^*, t \in \Sigma_-^*} r_T((s, t))$ can be computed with the formula

$$r_T(\Sigma^*) = \alpha_1^\top [I_d - M]^{-1} \alpha_\infty$$

where $M = \sum_{x \in \Sigma_+ \cup \{*\}, y \in \Sigma_- \cup \{*\}} M_{\frac{x}{y}}$.

Example 1. Let us consider a particular subclass of FST: $\Sigma_- = \Sigma_+ = \{0, 1\}$, where $M_0 = M_1 = M_{\frac{0}{*}} = M_{\frac{1}{*}} = 0$.



$$\alpha_1 = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \quad M_{\frac{*}{*}} = \begin{pmatrix} 1/3 & 0 \\ 0 & 1/6 \end{pmatrix} \quad M_{\frac{0}{*}} = \begin{pmatrix} 1/6 & 0 \\ 0 & 1/3 \end{pmatrix}$$

$$\alpha_\infty = \begin{pmatrix} 1/4 \\ 0 \end{pmatrix} \quad M_{\frac{0}{0}} = \begin{pmatrix} 0 & 1/4 \\ 0 & 0 \end{pmatrix} \quad M_{\frac{1}{1}} = \begin{pmatrix} 0 & 0 \\ 1/2 & 0 \end{pmatrix}$$

The FST A satisfies the constraints.

Let us draw all the paths for the i/o sequence $(01, 001)$. The dashed red edges are discarded because of the constraints. Thus, there are only two different non-zero paths, corresponding to $\frac{0*1}{001}$ (in green) and $\frac{*01}{001}$ (in blue).

Let us consider the model A which satisfies the constraints above. One has $r_A(\frac{0*1}{001}) = 1/96$, $r_A(\frac{*01}{001}) = 1/192$ and, as those two aligned sequences are the only possible alignments for $(01, 001)$, one has $r_A((01, 001)) = 1/64$. It is possible to check that $r_A(\Sigma^*) = 1$, thus the model computes a probability distribution.

2.3 Hankel Matrices

Let us recall some basic definitions and properties.

Let Σ be an alphabet, $U \subset \Sigma^*$ a set of prefixes, $V \subset \Sigma^*$ a set of suffixes. U is said to be *prefix-closed* if $uv \in U \Rightarrow u \in U$. V is said to be *suffix-closed* if $uv \in V \Rightarrow v \in V$.

Let us denote $U\Sigma$ the set $U \cup \{u_y^x | u \in U, x_y \in \Sigma\}$. Let us denote ΣV the set $V \cup \{x_y^v | v \in V, x_y \in \Sigma\}$.

A Hankel matrix on U and V is a matrix with rows corresponding to elements $u \in U$ and columns corresponding to $v \in V$, which satisfies $uv = u'v' \Rightarrow H(u, v) = H(u', v')$.

Definition 7. Let H a Hankel matrix for $U\Sigma$ and ΣV . One supposes that $\varepsilon \in U$ and $\varepsilon \in V$. One then defines the partial Hankel matrices defined for $u \in U$ and $\varepsilon \in V$:

$$H_\varepsilon(u, v) = H(u, v) \quad , \quad H_y^x(u, v) = H(u, \frac{x}{y}v) \quad , \quad H_1(v) = H(\varepsilon, v) \quad , \quad H_\infty(u) = H(u, \varepsilon)$$

The main result that we will be using is the following:

Proposition 2. *Let H a Hankel matrix for $U\Sigma$ and ΣV . One supposes that U is prefix-closed, V is suffix-closed, and that $\text{rank}(H_\varepsilon) = \text{rank}(H)$. Then the WA defined by*

$$\alpha_1^\top = H_1^\top H_\varepsilon^+ \quad , \quad \alpha_\infty = H_\infty \quad , \quad M_y = H_y H_\varepsilon^+$$

computes a mapping f such that $\forall u \in U, \forall v \in V, f(uv) = H_\varepsilon(u, v)$.

We will not give a proof of this result, as a more general result is seen further. The rank equality comes from the fact that the WA defined above has the same rank as H_ε , and that the rank of a mapping f which satisfies $f(uv) = H(u, v)$ is at least the rank of H . The following example shows that the prefix and suffix closeness are necessary conditions.

Example 2. *Let us consider the following Hankel over the set of prefixes $U\Sigma$ and the set of suffixes ΣV with $U = \{\varepsilon, a^3, b^3\}$, and $V = \{\varepsilon, a^3\}$.*

$$H_\varepsilon = \begin{matrix} & \varepsilon & a^3 \\ \varepsilon & 0 & 0 \\ a^3 & 0 & 1/4 \\ b^3 & 0 & 1/4 \end{matrix}, \quad H = \begin{matrix} & \varepsilon & a & a^3 & a^4 \\ \varepsilon & 0 & 0 & 0 & 0 \\ a & 0 & 0 & 0 & 0 \\ a^3 & 0 & 0 & 1/4 & 1/4 \\ b^3 & 0 & 0 & 1/4 & 1/4 \\ a^4 & 0 & 0 & 1/4 & 1/4 \end{matrix}, \quad H' = \begin{matrix} & a^2 & a^3 & a^4 \\ a^2 & 0 & 0 & 1/4 \\ b^2 & 0 & 1/4 & 1/4 \\ b^3 & 1/4 & 1/4 & 1/4 \\ b^4 & 1/4 & 1/4 & 1/4 \end{matrix}$$

One has $\varepsilon \in U$ and $\varepsilon \in V$, and also $\text{rank}(H_\varepsilon) = \text{rank}(H) = 1$, thus the computed WA is rank 1. Such a WA cannot compute a mapping such that $r_T(\varepsilon) = 0$ and $r_T(a^6) = 1/4$. The complete Hankel matrix has at least rank 7.

3 Inducing FST from Generalized Hankel Matrices

Proposition 2 tells us that if we had access to certain sub-blocks of the Hankel matrix for aligned sequences we could recover the FST model. However, we do not have access to the hidden alignment information: we only have access to the statistics $p(s, t)$, which we will call *observations*. One natural idea would be to search for a Hankel matrix that agrees with the observations. To do so, we introduce *observable constraints*, which are linear constraints of the form $p(s, t) = \sum_{y_1 \dots y_n \in [s, t]} r_T(x_1 \dots x_n, y_1 \dots y_n)$, where $r_T(x_1 \dots x_n, y_1 \dots y_n)$ is computed from the Hankel.

From a matrix satisfying the hypothesis of Proposition 2 and the observation constraints, one can derive an FST computing a mapping which agrees on the observations.

Given an i/o sequence (s, t) , the size of $[s, t]$ (hence the size of the Hankel matrix) is in general exponential in the size of s and t . In order to overcome that problem when considering the observation constraints, one will consider aggregations of rows and columns, corresponding to sets of prefixes and suffixes. Obviously, the definition of a Hankel matrix must be extended to this case.

One will denote by $\mathcal{P}(A)$ the set of subsets of A .

Definition 8. *Let $u, u' \in \mathcal{P}(\Sigma^*)$. The set uu' is defined by $uu' = \{ww' \mid w \in u, w' \in u'\}$.*

We will denote sets of alignments as follows: $\begin{smallmatrix} x_{1:n} \\ y_{1:n} \end{smallmatrix}$ will denote the set $\{x_{1:n}, y_{1:n}\}$, which contains a single aligned sequence; then $\begin{smallmatrix} x_{1:n} \\ y_{1:n} \end{smallmatrix} [s, t]$ will denote the set $\{x_{1:n}x_{n+1:n+k}, y_{1:n}y_{n+1:n+k} \mid x_{n+1:n+k} \in [s, t]\}$, which extends $\{x_{1:n}, y_{1:n}\}$ with all ways of aligning (s, t) . We will also use $[s, t]_{y_{1:n}}^{x_{1:n}}$ analogously.

Definition 9. *A generalized prefix (resp. generalized suffix) is the empty set or a set of the form $[s, t]_{y_{1:n}}^{x_{1:n}}$ (resp. $_{y_{1:n}}^{x_{1:n}} [s, t]$), where the aligned part is possibly empty.*

3.1 Generalized Hankel

One needs to extend the definition of a Hankel matrix to the generalized prefixes and suffixes.

Definition 10. *Let U be a set of generalized prefixes, V be a set of generalized suffixes. Let H be a matrix indexed by U (in rows) and V (in columns). H is a generalized Hankel matrix if it satisfies:*

$$\forall \bar{u}, \bar{u}' \subset U, \forall \bar{v}, \bar{v}' \subset V, \quad \bigcup_{u \in \bar{u}, v \in \bar{v}} uv = \bigcup_{u' \in \bar{u}', v' \in \bar{v}'} u'v' \Rightarrow \sum_{u \in \bar{u}, v \in \bar{v}} H(u, v) = \sum_{u' \in \bar{u}', v' \in \bar{v}'} H(u', v')$$

where \bigcup is the disjoint union.

In particular, if U and V are sets of regular prefixes and suffixes, this definition encompasses the regular definition for a Hankel matrix.

Definition 11. Let U be a set of generalized prefixes. U is prefix-closed if

$$\begin{aligned} [s, t]_{y_{1:n}}^{x_{1:n}} \in U &\Rightarrow [s, t]_{y_{1:n-1}}^{x_{1:n-1}} \in U \\ [s_{1:n}, t_{1:k}] \in U &\Rightarrow [s_{1:n-1}, t_{1:k}]_*^{s_n}, [s_{1:n}, t_{1:k-1}]_{t_k}^*, [s_{1:n-1}, t_{1:k-1}]_{t_k}^{s_n} \in U \end{aligned}$$

Definition 12. Let V be a set of generalized suffixes. V is suffix-closed if

$$\begin{aligned} [s, t]_{y_{1:n}}^{x_{1:n}} \in V &\Rightarrow [s, t]_{y_{2:n}}^{x_{2:n}} \in V \\ [s_{1:n}, t_{1:k}] \in V &\Rightarrow [s_{2:n}, t_{1:k}]_*^{s_1}, [s_{1:n}, t_{2:k}]_{t_1}^*, [s_{2:n}, t_{2:k}]_{t_1}^{s_1} \in V \end{aligned}$$

Definition 13. Let U be a set of generalized prefixes, V be a set of generalized suffixes. The right-operator completion of U is the set $U\Sigma = U \cup \{u_y^x | u \in U, x_y \in \Sigma\}$. The left-operator completion of V is the set $\Sigma V = V \cup \{x_y^v | v \in V, x_y \in \Sigma\}$.

A key result is the following, which is analogous to Proposition 2 for generalized Hankel matrices:

Proposition 3. Let U and V be two sets of generalized prefixes and generalized suffixes. Let H be a Hankel matrix built from $U\Sigma$ and ΣV . One supposes that $\text{rank}(H_\varepsilon) = \text{rank}(H)$, U is prefix-closed and V is suffix-closed. Then the model A defined by $\alpha_1^\top = H_1^\top H_\varepsilon^+$, $\alpha_\infty = H_\infty$, $M_x = H_y^x H_\varepsilon^+$ computes a mapping r_A such that

$$\forall u \in U, \forall v \in V, r_A(uv) = H_\varepsilon(u, v)$$

Proof. The proof can be found in the Appendix. □

Let S be a sample of unaligned sequences. Let $\text{pref}_{S,in}$ (resp. $\text{pref}_{S,out}$, $\text{suff}_{S,in}$, $\text{suff}_{S,out}$) be the prefix (resp. suffix) closure of input (resp. output) strings in S . Let $U = \{[s, t]\}_{s \in \text{pref}_{S,in}, t \in \text{pref}_{S,out}}$ and $V = \{[s, t]\}_{s \in \text{suff}_{S,in}, t \in \text{suff}_{S,out}}$. The sets U and V contain all the observed pairs of unaligned sequences, and one can check that the sizes of $U\Sigma$ and ΣV are polynomial in the size of S .

Example 3. Let us continue with the same model A as in Example 1. Let us now consider the prefixes $\epsilon, \frac{0}{1}$ and the suffixes $\epsilon, \frac{1}{1}$. The Hankel matrices will be:

$$\begin{aligned} H_1 &= \begin{pmatrix} 1/4 \\ 0 \end{pmatrix} H_\infty = \begin{pmatrix} 1/4 \\ 0 \end{pmatrix} H_\varepsilon = \begin{pmatrix} 1/4 & 0 \\ 0 & 1/32 \end{pmatrix} H_0^* = \begin{pmatrix} 1/24 & 0 \\ 0 & 1/96 \end{pmatrix} \\ H_1^* &= \begin{pmatrix} 1/12 & 0 \\ 0 & 1/192 \end{pmatrix} H_0 = \begin{pmatrix} 0 & 1/32 \\ 0 & 0 \end{pmatrix} H_1 = \begin{pmatrix} 0 & 0 \\ 1/32 & 0 \end{pmatrix} \end{aligned}$$

and one finally gets the model A' defined by:

$$\begin{aligned} \alpha_1 &= \begin{pmatrix} 1 \\ 0 \end{pmatrix} \alpha_\infty = \begin{pmatrix} 1/4 \\ 0 \end{pmatrix} M_1^* = \begin{pmatrix} 1/3 & 0 \\ 0 & 1/6 \end{pmatrix} \\ M_0^* &= \begin{pmatrix} 1/6 & 0 \\ 0 & 1/3 \end{pmatrix} M_0 = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix} M_1 = \begin{pmatrix} 0 & 0 \\ 1/8 & 0 \end{pmatrix} \end{aligned}$$

One can easily check that A' computes the same probability distribution than A .

4 FST Learning as Non-convex Optimization

Proposition 3 shows that FST models can be recovered from generalized Hankel matrices. In this section we show how FST learning can be framed as an optimization problem where one searches for a low-rank generalized Hankel matrix that agrees with observation constraints derived from a sample. We assume here that p is a probability distribution over i/o sequences.

We will denote by $z = (p([s, t]))_{s \in \Sigma_+^*, t \in \Sigma_-^*}$ the vector built from observable probabilities, and by $z_S = (p_S([s, t]))_{s \in \Sigma_+^*, t \in \Sigma_-^*}$ the set of empirical observable probabilities, where p_S is the frequency deduced from an i.i.d. sample S .

Let \vec{H} be the vector describing the coefficients of H . Let K be the matrix such that $K\vec{H} = 0$ represents the Hankel constraints (cf. definition 10). Let O be the matrix such that $O\vec{H} = \mathbf{z}$ represents the observation constraints (i.e. $\sum H([s, t], \epsilon) = p([s, t])$).

The optimization task is the following:

$$\begin{aligned} & \underset{H}{\text{minimize}} \quad \text{rank}(H) \\ & \text{subject to} \quad \|O\vec{H} - \mathbf{z}\|_2 \leq \mu \\ & \quad \quad \quad K\vec{H} = \mathbf{0} \\ & \quad \quad \quad \|H\|_2 \leq 1. \end{aligned} \tag{1}$$

The condition $\|H\|_2 \leq 1$ is necessary to bound the set of possible solutions. In particular, if H is the Hankel matrix of a probability distribution, the condition is satisfied: one has $\|H\|_1 \leq 1$ as each column is a probability distribution, as well as $\|H\|_F \leq 1$ and thus $\|H\|_2 \leq \sqrt{\|H\|_1 \|H\|_F} \leq 1$. Let us remark that the set of matrices satisfying the constraints is a compact set.

Let us denote \mathcal{H}_μ the class of Hankel matrices solutions of (1) for a given μ , where \mathbf{z} represents the $p(u_i v_j)$. Let us denote \mathcal{H}_μ^S the class of Hankel matrices solutions of (1) for a given μ , where \mathbf{z}_S represents $p_S(u_i v_j)$, the observed frequencies in a sample S i.i.d. with respect to p .

Proposition 4. *Let p be a distribution over i/o sequences computed by an FST. There exists U and V such that any solution of \mathcal{H}_0 leads to an FST*

$$\alpha_1^\top = H_1^\top H_\varepsilon^+, \alpha_\infty = H_\infty, M_y^x = H_y^x H_\varepsilon^+$$

which computes p .

Proof. The proof can be found in the Appendix. \square

4.1 Theoretical Properties

We now present the main theoretical results related to the optimization problem (1). The first one concerns the rank identification, while the second one concerns the consistency of the method.

Proposition 5. *Let p be a rank d distribution computed by an FST. There exists μ_2 such that for any $\delta > 0$ and any i.i.d. sample S*

$$|S| > \left(\frac{2 + \sqrt{8 \log(1/\delta)}}{\mu_2} \right)^2$$

implies that any $H \in \mathcal{H}_\mu^S$ solution of (1) with $\mu = \frac{1 + \sqrt{2 \log(1/\delta)}}{\sqrt{|S|}}$ leads to a rank- d FST with probability at least $1 - \delta$.

Proof. The proof can be found in the Appendix. \square

Proposition 6. *Let p be a rank d distribution computed by an FST. There exists μ_ϵ such that for any $\delta > 0$ and any i.i.d. sample S*

$$|S| > \left(\frac{1 + \sqrt{2 \log(1/\delta)}}{\mu_\epsilon} \right)^2$$

implies that for any $H \in \mathcal{H}_\mu^S$ solution of (1) with $\mu = \frac{1 + \sqrt{2 \log(1/\delta)}}{\sqrt{|S|}}$ leading to a model A_s there exists a model A computing p such that

$$|A, A_S|_\infty \leq O\left(\frac{d^2 \epsilon}{\sigma_p^3}\right)$$

where $|A, A_S|_\infty$ is the maximum distance for model parameters, and σ_p is a non-zero parameter depending on p .

Proof. The proof can be found in the Appendix. \square

Example 4. This continues Examples 1 and 3. Let us first remark that, among all the values used to build the Hankel matrices in Example 3, some of them correspond to observable statistics, as there is only one possible alignment for them. Conversely, the exact values of $r_M(\begin{smallmatrix} 0 & * & 1 \\ 0 & 0 & 1 \end{smallmatrix})$ and $r_M(\begin{smallmatrix} 0 & 1 & 1 \\ 0 & * & 1 \end{smallmatrix})$ are not observable. Then the rank minimization objective is not sufficient, as it allows any value for those variables.

Let us consider now larger sets of prefixes and suffixes: $\varepsilon, \begin{smallmatrix} * \\ 0 \end{smallmatrix}, \begin{smallmatrix} 0 \\ 0 \end{smallmatrix}$ and $\varepsilon, \begin{smallmatrix} 1 \\ * \end{smallmatrix}, \begin{smallmatrix} 1 \\ 1 \end{smallmatrix}$. One then has

$$H_\varepsilon = \begin{pmatrix} 1/4 & 1/12 & 0 \\ 1/24 & ? & 0 \\ 0 & 0 & 1/32 \end{pmatrix} H_1 = \begin{pmatrix} 1/12 & 1/36 & 0 \\ ? & ? & 0 \\ 0 & 0 & ? \end{pmatrix} H_1 = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1/32 & ? & 0 \end{pmatrix}$$

We want to minimize the rank of H' subject to the constraints O and K :

$$H' = \begin{pmatrix} H_\varepsilon \\ H_1^* \\ H_1^1 \end{pmatrix} = (h_{ij}), \quad O = \begin{cases} h_{11} = 1/4 & \vdots \\ h_{12} = 1/12 & h_{63} + h_{92} = 1/64 \\ h_{13} = 0 & \vdots \\ h_{21} = 1/24 & \end{cases}, \quad K = \begin{cases} h_{12} = h_{41} & h_{23} = h_{81} \\ h_{13} = h_{71} & h_{32} = h_{61} \\ h_{22} = h_{51} & h_{33} = h_{91} \end{cases}$$

The relation $h_{63} + h_{92} = 1/64$ is due to fact that $r_A(\begin{smallmatrix} 0 & * & 1 \\ 0 & 0 & 1 \end{smallmatrix}) + r_A(\begin{smallmatrix} * & 0 & 1 \\ 0 & 0 & 1 \end{smallmatrix}) = r_A((01, 001)) = 1/64$. One has $h_{22} = h_{51}$ as they both represent $p(\begin{smallmatrix} * & 1 \\ 0 & * \end{smallmatrix})$. It is clear that H' has a rank greater or equal than 2. The only way to reach rank 2 under the constraints is

$$h_{22} = h_{51} = 1/72, h_{52} = 1/216, h_{63} = 1/192, h_{92} = 1/96$$

Thus, the process of rank minimization under linear constraints leads to one single model, which is identical to the original one. Of course, in the general case, the rank minimization objective may lead to several models.

4.2 Convex Relaxation

The problem as it is stated in (1) is NP-hard to solve in the size of the Hankel matrix, hence impossible to deal with in practical cases. One can solve instead a convex relaxation of the problem (1), obtained by replacing the rank objective by the nuclear norm. The relaxed optimization statement is then the following:

$$\begin{aligned} & \underset{H}{\text{minimize}} && \|H\|_* \\ & \text{subject to} && \|O\tilde{H} - \mathbf{z}\|_2 \leq \mu \\ & && K\tilde{H} = \mathbf{0} \\ & && \|H\|_2 \leq 1. \end{aligned} \tag{2}$$

This type of relaxation has been used extensively in multiple settings [13]. The nuclear norm $\|\cdot\|_*$ plays the same role than the $\|\cdot\|_1$ norm in convex relaxations of the $\|\cdot\|_0$ norm, used to reach sparsity under linear constraints.

5 Experiments

We ran synthetic experiments using samples generated from random FST with input-output alphabets of size two. The main goal of our experiments was to compare our algorithm to a *supervised* spectral algorithm for FST that has access to alignments. In both methods, the Hankel was defined for prefixes and suffixes up to length 1. Each run consists of generating a target FST, and generating N aligned samples according to the target distribution. These samples were directly used by the supervised algorithm. Then, we removed the alignment information from each sample (thus obtaining a pair of unaligned strings), and we used them to train an FST with our general learning algorithm, trying different values for a C parameter that trades-off the nuclear norm term and the observation term. We ran this experiment for 8 target models of 5 states, for different sampling sizes. We measure the L_1 error of the learned models with respect to the target distribution on all unaligned pairs of strings whose sizes sum up to 6. We report results for geometric averages.

In addition, we ran two additional methods. First a *factorized* method that assumes that the two sequences are generated independently, and learns two separate weighted automata using a spectral

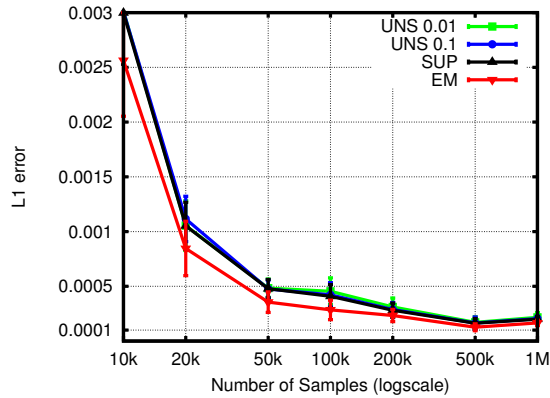


Figure 1: Learning curves for different methods: SUP, supervised; UNS, unsupervised with different regularizers; EM, expectation maximization. The curves are averages of L1 error for random target models of 5 states.

method. Its performance is very bad, with L1 error rates at ~ 0.08 , which confirms that our target models generate highly dependent sequence pairs. This baseline result also implies that the rest of the methods can learn the dependencies between paired strings. Second, we ran an *Expectation Maximization* algorithm (EM).

Figure 1 shows the performance of the learned models with respect to the number of samples. Clearly, our algorithm is able to estimate the target distribution and gets close to the performance of the supervised method, while making use of much simpler statistics. EM performed slightly better than the spectral methods, but nonetheless at the same levels of performance.

One can find other experimental results for the unsupervised spectral method in [1], where it is shown that, under certain circumstances, unsupervised spectral method can perform better than supervised EM. Though the framework (unsupervised learning of PCFGs) is not the same, the method is similar and the optimization statement is identical.

6 Conclusion

In this paper we presented a spectral algorithm for learning FST from unaligned sequences. This is the first paper to derive a spectral algorithm for the unsupervised FST learning setting. We prove that there is theoretical identifiability of the rank and the parameters of an FST distribution, using a rank minimization formulation. However, this problem is NP-hard, and it remains open whether there exists a polynomial method with identifiability results. Classically, rank minimization problems are solved with convex relaxations such as the nuclear norm minimization we have proposed. In experiments, we show that our method is comparable to a fully supervised spectral method, and close to the performance of EM.

Our approach follows a similar idea to that of [3] since both works combine classic ideas from spectral learning with classic ideas from low rank matrix completion. The basic idea is to frame learning of distributions over structured objects as a low-rank matrix factorization subject to linear constraints derived from observable statistics. This method applies to other grammatical inference domains, such as unsupervised spectral learning of PCFGs ([1]).

Acknowledgments

We are grateful to Borja Balle and the anonymous reviewers for providing us with helpful comments. This work was supported by a Google Research Award, and by projects XLike (FP7-288342), ERA-Net CHISTERA VISEN, TACARDI (TIN2012-38523-C02-02), BASMATI (TIN2011-27479-C04-03), and SGR-LARCA (2009-SGR-1428). Xavier Carreras was supported by the Ramón y Cajal program of the Spanish Government (RYC-2008-02223).

References

- [1] R. Bailly, X. Carreras, F. M. Luque, and A. Quattoni. Unsupervised spectral learning of wcfg as low-rank matrix completion. *EMNLP*, 2013.
- [2] R. Bailly, F. Denis, and L. Ralaivola. Grammatical inference as a principal component analysis problem. In *Proc. ICML*, 2009.
- [3] B. Balle and M. Mohri. Spectral learning of general weighted automata via constrained matrix completion. In *Proc. of NIPS*, 2012.
- [4] B. Balle, A. Quattoni, and X. Carreras. A spectral learning algorithm for finite state transducers. *ECML-PKDD*, 2011.
- [5] Borja Balle, Ariadna Quattoni, and Xavier Carreras. Local loss optimization in operator models: A new insight into spectral learning. In John Langford and Joelle Pineau, editors, *Proceedings of the 29th International Conference on Machine Learning (ICML-12)*, ICML '12, pages 1879–1886, New York, NY, USA, July 2012. Omnipress.
- [6] M. Bernard, J-C. Janodet, and M. Sebban. A discriminative model of stochastic edit distance in the form of a conditional transducer. *Grammatical Inference: Algorithms and Applications*, 4201, 2006.
- [7] B. Boots, S. Siddiqi, and G. Gordon. Closing the learning planning loop with predictive state representations. *I. J. Robotic Research*, 2011.
- [8] F. Casacuberta. Inference of finite-state transducers by using regular grammars and morphisms. *Grammatical Inference: Algorithms and Applications*, 1891, 2000.
- [9] A. Clark. Partially supervised learning of morphology with stochastic transducers. In *Proc. of NLPRS*, pages 341–348, 2001.
- [10] Shay B. Cohen, Karl Stratos, Michael Collins, Dean P. Foster, and Lyle Ungar. Spectral learning of latent-variable pcfgs. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 223–231, Jeju Island, Korea, July 2012. Association for Computational Linguistics.
- [11] J. Eisner. Parameter estimation for probabilistic finite-state transducers. In *Proc. of ACL*, pages 1–8, 2002.
- [12] G. Stewart et J.-G. Sun. Matrix perturbation theory. *Academic Press*, 1990.
- [13] Maryam Fazel. *Matrix rank minimization with applications*. PhD thesis, Stanford University, Electrical Engineering Dept., 2002.
- [14] D. Hsu, S. M. Kakade, and T. Zhang. A spectral algorithm for learning hidden markov models. In *Proc. of COLT*, 2009.
- [15] Mehryar Mohri. Finite-state transducers in language and speech processing. *Computational Linguistics*, 23(2):269–311, 1997.
- [16] A.P. Parikh, L. Song, and E.P. Xing. A spectral algorithm for latent tree graphical models. *ICML*, 2011.
- [17] S.M. Siddiqi, B. Boots, and G.J. Gordon. Reduced-rank hidden markov models. *AISTATS*, 2010.
- [18] L. Song, B. Boots, S. Siddiqi, G. Gordon, and A. Smola. Hilbert space embeddings of hidden markov models. *ICML*, 2010.
- [19] L. Zwald and G. Blanchard. On the convergence of eigenspaces in kernel principal component analysis. *NIPS*, 2005.

7 Appendix

7.1 Proof of the Proposition 3

In order to prove that result, one needs some intermediate results. Let H_U (resp. $H_{U_y^x}$) be the submatrix of H corresponding to prefixes in U (resp. of the form u_y^x with $u \in U$). Let H_V (resp. H_{yV}^x) be the submatrix of H corresponding to suffixes in V (resp. of the form y_v^x with $v \in V$).

Lemma 1. *Let u and v be two vectors such that $u^\top H_\varepsilon = v^\top H_\varepsilon$. Then, for $x \in \Sigma_+ \cup \{*\}$, $y \in \Sigma_- \cup \{*\}$, one has $u^\top H_y^x = v^\top H_y^x$.*

Proof. H_ε is a submatrix of H_U with the same rank.

Let u and v be two vectors such that $u^\top H_\varepsilon = v^\top H_\varepsilon$, then $u^\top H_U = v^\top H_U$ because H_ε and H_U have the same rank. Thus, as each H_y^x is a submatrix of H_U , one has $u^\top H_y^x = v^\top H_y^x$. \square

Lemma 2. *Let $u \in U\Sigma$. Then the vector*

$$\sum_{\substack{x_1 \dots x_n \\ y_1 \dots y_n \in u}} (H_{y_n}^{x_n})^\top \dots (H_{y_1}^{x_1} H_\varepsilon^+)^\top ((H_\varepsilon^+)^\top H_1)$$

is the row of H_V corresponding to the prefix u . In particular, if $u \in U$, the vector is equal to the row of H_ε corresponding to the prefix u .

Proof. By induction. H_1 is the row of H_V corresponding to ε .

1) Let us suppose that $u = u_y^x$. Because $U\Sigma$ is prefix-closed, one has $u' \in U\Sigma$. Let z' be the row of H_V corresponding to u' . $(H_\varepsilon^+)^\top z'$ represents a decomposition of z' in terms of rows of H_ε . The vector $(H_y^x)^\top (H_\varepsilon^+)^\top z$ is the same linear combination of rows of H_{yV}^x , and by rank equality is the same as the row of H_{yV}^x corresponding to u' . Because H is a Hankel matrix, it is equal to the row of H_V corresponding to $u_y^x = u$.

2) Let us suppose that $u = [s_{1:n}, t_{1:k}]$. Then $u_1 = [s_{1:n-1}, t_{1:k}]_{*}^{s_n} \in U\Sigma$, $u_2 = [s_{1:n}, t_{1:k-1}]_{t_k}^* \in U\Sigma$, $u_3 = [s_{1:n-1}, t_{1:k-1}]_{t_k}^{s_n} \in U\Sigma$. With the same argument as before applied to u_1 , u_2 and u_3 , and because H is Hankel, one has the result. \square

One has then the symmetric result for the suffixes.

Lemma 3. *Let $v \in \Sigma V$. Then the vector*

$$\sum_{\substack{x_1 \dots x_n \\ y_1 \dots y_n \in v}} (H_{y_1}^{x_1}) \dots (H_{y_n}^{x_n} H_\varepsilon^+)(H_\infty)$$

is the column of H_U corresponding to the suffix v . In particular, if $v \in V$, the vector is equal to the column of H_ε corresponding to the suffix v .

Proof. It is just the symmetric case of the previous lemma. \square

7.1.1 Proof of the Proposition 3

Let $u \in U$, $v \in V$. Let H_u be the row of H_ε corresponding to u , H_v the column of H_ε corresponding to v . One then has, by Lemma 2 and Lemma 3, $r_M(uv) = H_u^\top H_\varepsilon^+ H_v$. The vector $H_\varepsilon^+ H_v$ represents a decomposition of H_v equivalent to the vector $\mathbf{1}_v$. Then $r_M(uv) = H_u^\top \mathbf{1}_v = H_\varepsilon(u, v)$. \square

7.2 Proof of the Proposition 4

Definition 14. *Let p be a distribution over i/o sequences computed by an FST. Let $\text{rank}(p)$ be the minimal integer d such that there exist an FST with d states computing p . Let \mathcal{V}_p be the class of parameters for all rank- d FSTs over bi-sequences which compute the same distribution over i/o sequences as p .*

Definition 15. An affine variety is the set of solutions of a (maybe infinite) polynomial equation system:

$$\begin{cases} P_1(X_1, \dots, X_n) = 0 \\ \vdots \end{cases}$$

Lemma 4. Let p be a rank d distribution over bi-sequences computed by an FST. Then \mathcal{V}_p is an affine variety.

Proof. Let A be a d -state FST. The value computed by A for a given i/o sequence (s, t) is a polynomial in its parameter denoted $P_{(s,t)}$. Thus, the set of parameters corresponding to d -state FST computing a given value $p((s, t))$ for (s, t) is an affine variety defined by $\{(X_1, \dots, X_n) | P_{(s,t)} - p((s, t)) = 0\}$, and \mathcal{V}_p is the affine variety defined by: $\bigcap_{(s_i, t_j) \in \Sigma_+ \times \Sigma_-} \{(X_1, \dots, X_n) | P_{(s_i, t_j)} - p((s_i, t_j)) = 0\}$. \square

Lemma 5. Let p be a rank d distribution over bi-sequences computed by an FST. Then there exists a finite set G_p of i/o sequences, such that $\mathcal{V}_p = \bigcap_{(s_i, t_j) \in G_p} \{(X_1, \dots, X_n) | P_{(s_i, t_j)} - p((s_i, t_j)) = 0\}$. Such a set G_p is called a generative set for p .

Proof. The ring $\mathbb{R}[X_1, \dots, X_n]$ is Noetherian, in particular the sequence $I_k = \bigcap_{k' \leq k} \{(X_1, \dots, X_n) | P_{(s_{i_{k'}}, t_{j_{k'}})}(X_1, \dots, X_n) - p((s_{i_{k'}}, t_{j_{k'}})) = 0\}$ is stationary. One has $\mathcal{V}_p = \bigcup_n I_n = \bigcup_{n \leq N} I_n$ for a certain N . One can take $G_p = \bigcup_{n \leq N} (s_{i_n}, t_{j_n})$. \square

Corollary 1. Let p be a rank d distribution over i/o sequences computed by an FST. Let G_p be a generative set for p . Let A be an FST of rank $\leq d$. One then has:

$$r_A|_{G_p} = p|_{G_p} \Leftrightarrow r_A = p$$

7.2.1 Proof of Proposition 4

Proof. Let p be a rank d distribution over i/o sequences computed by an FST. Let G_p be a generative set for p . Let U_0 (resp. V_0) be the prefix-closure (resp. suffix-closure) of G_p . Let $U_{i+1} = U_i \Sigma$, $U = U_{d+1}$ and $V_{i+1} = \Sigma V_i$, $V = V_{d+1}$. Let H_i be the minimum rank Hankel matrix over U_i and V_i , and let H be a minimum rank Hankel matrix over U and V . With Corollary 1 and Proposition 3, it is sufficient to prove that $\text{rank}(H_d) = \text{rank}(H) = d$. As the Hankel matrix of p fulfills the hypothesis, one has $\text{rank}(H) \leq d$. Among the family of $(d+1)$ couples $(H_0, H_1), \dots, (H_d, H)$, one of them satisfies $\text{rank}(H_i) = \text{rank}(H_{i+1})$, because otherwise $\text{rank}(H_i)$ would take $d+2$ different values between 0 and d . Thus, the FST computed from H_{i+1} agrees on G_p with p by Proposition 3, and by Corollary 1, as $G_p \subset U \times V$, this FST computes p . By minimality of the rank, one has $\text{rank}(H_i) = \text{rank}(H_{i+1}) = d$, and thus $\text{rank}(H_d) = \text{rank}(H) = d$. \square

7.3 Proof of the Proposition 5

Lemma 6. Let p be a rank d distribution computed by an FST. Let U and V be such as in Proposition 4. There exists $\sigma > 0$ such $H \in \mathcal{H}_0 \Rightarrow \sigma_d(H_\varepsilon) \geq \sigma$, where $\sigma_d(H_\varepsilon)$ is the d -th singular value of H_ε .

Proof. For $\mu = 0$, the rank minimization is equivalent to $\text{rank}(H) \leq d$, thus the set \mathcal{H}_0 of the solutions of (1) is a closed bounded set, thus compact. Suppose that the assumption is false, this means, by compactity, that one can find a sequence H_n such that $\sigma_d(H_{n\varepsilon})$ converges towards a matrix H_ω such that $\sigma_d(H_{\omega\varepsilon}) = 0$ by continuity of singular values. As $H_\omega \in \mathcal{H}_0$, The FST obtained from H_ω computes p , which contradicts the fact that $\text{rank}(H_{\omega\varepsilon}) = d$ (cf. proof of Proposition 4). \square

Lemma 7. Let p be a distribution computed by a rank d FST. Let U and V be such as in Proposition 4. Let σ be as in Lemma 6. There exists μ_2 such that $H \in \mathcal{H}_{\mu_2} \Rightarrow \sigma_d(H_\varepsilon) > \sigma/2$.

Proof. Suppose the assumption is false: there exists a convergent sequence of Hankel matrices $H_n \in \mathcal{H}_{1/n}$ such that $\sigma_d(H_{n\varepsilon}) < \sigma/2$, and whose limit is M_ω . One then has $H_\omega \in \mathcal{H}_0$, and $\sigma_d(H_{\omega\varepsilon}) \leq \sigma/2$ by continuity, which contradicts Lemma 6. \square

In particular, this implies that, for a certain μ_2 , all the solutions \mathcal{H}_{μ_2} of (1) will be such that H_ε is rank d , thus \mathcal{H}_{μ_2} is compact.

Lemma 8. *Let p be a rank d distribution computed by an FST. Let U and V be such as in Proposition 6. For all $\epsilon > 0$ there exists μ_ϵ such that $H \in \mathcal{H}_{\mu_\epsilon} \Rightarrow \min_{H_0 \in \mathcal{H}_0} (\|H - H_0\|_F) \leq \epsilon$.*

Proof. Let us consider $\mu_\epsilon < \mu_2$, μ_2 being as in Lemma 7. The rank minimization is equivalent to $\text{rank}(H) \leq d$, thus the set $\mathcal{H}_{\mu_\epsilon}$ is compact. Let us suppose that the assumption is false, and that there exists a sequence \mathcal{H}_n such that $H_n \in \mathcal{H}_{1/n}$ and $\min_{H_0 \in \mathcal{H}_0} (\|H_n - H_0\|_F) > \epsilon$. The limit H_ω belongs to \mathcal{H}_0 and satisfies $\min_{H_0 \in \mathcal{H}_0} (\|H_\omega - H_0\|_F) \geq \epsilon$ which is contradictory. \square

Lemma 9. *Let p be a rank d distribution computed by an FST. Let U and V be such as in Proposition 4. Let $\delta > 0$ be a confidence parameter. Let S be an i.i.d. sample of size N , drawn with respect to p . Let $z_S = (p_S([s, t]))_{[s, t] \in U}$ be the vector of frequencies in the sample S , and let $z = (p([s, t]))_{[s, t] \in U}$. One has, with probability at least $1 - \delta$:*

$$\|z - z_S\|_2 < \frac{1 + \sqrt{2 \log(1/\delta)}}{\sqrt{N}}$$

Proof. Let S_i be a sample differing from S for the i -th entry. One has $\|z_S - z_{S_i}\|_2 \leq \sqrt{2}/N = c_i$. One also has $\mathbb{E}(\|z - z_S\|_2^2) \leq 1/N$ because of the variance of a multinomial, and thus $\mathbb{E}(\|z - z_S\|_2) \leq \sqrt{\mathbb{E}(\|z - z_S\|_2^2)} \leq 1/\sqrt{N}$.

Applying the McDiarmid's inequality gives $\mathbb{P}(\|z_S - z\|_2 \geq \mathbb{E}(\|z - z_S\|_2) + \epsilon) \leq e^{-\frac{\epsilon^2}{2 \sum c_i^2}}$. With $\delta = e^{-\frac{\epsilon^2}{2 \sum c_i^2}} = e^{-\frac{N \epsilon^2}{4}}$, thus $\epsilon = \sqrt{\frac{2 \log(1/\delta)}{N}}$, one has the result. \square

7.3.1 Proof of the Proposition 5

Let μ_2 be as in Lemma 7. By the Lemma 9, with probability $1 - \delta$, one has $\mathcal{H}_0 \subset \mathcal{H}_\mu^S$, thus $\text{rank}(H) \leq d$ for any $H \in \mathcal{H}_\mu^S$. Moreover, as $\mathcal{H}_\mu^S \subset \mathcal{H}_{2\mu}$, the condition $\mu < \mu_2$ implies that $\text{rank}(H_\varepsilon) \geq d$ for any $H \in \mathcal{H}_\mu^S$. \square

7.4 Proof of the Proposition 6

Lemma 10. *Let p be a rank d distribution computed by an FST. Let S be an i.i.d. sample of size N with respect to p . Let $\delta > 0$ be a confidence parameter. For any $\epsilon > 0$, let μ_ϵ be as in Lemma 8. One supposes that*

$$N > \left(\frac{1 + \sqrt{2 \log(1/\delta)}}{\mu_\epsilon} \right)^2$$

With probability $1 - \delta$, for any $H \in \mathcal{H}_{\mu_\epsilon}^S$, $\min_{H_0 \in \mathcal{H}_0} (\|H - H_0\|_F) < \epsilon$.

Proof. This is just Lemma 8 and Lemma 9 together. \square

Let us define the distance between two models with the same rank:

Definition 16. *Let $A = (\alpha_1, \alpha_\infty, M_x)$ and $A' = (\alpha'_1, \alpha'_\infty, M'_x)$ be two FSTs with d states, on the same alphabet. One defines the distance*

$$|A, A'|_\infty = \max \left(\max_i (|(\alpha_1)_i - (\alpha'_1)_i|), \max_i (|(\alpha_\infty)_i - (\alpha'_\infty)_i|), \max_{i,j,x,y} (|(M_x)_{i,j} - (M'_x)_{i,j}|) \right)$$

Let us recall a result [12]:

Lemma 11. *Let H and $H' = H + E$ be two $n \times m$ matrices. Let $\sigma_1 \geq \dots \geq \sigma_n$ be the singular values of H , and let $\sigma'_1 \geq \dots \geq \sigma'_n$ be the singular values of H' . One then has*

$$|\sigma_i - \sigma'_i| \leq \|E\|_2$$

Let $H = L^\top D R$ and $H' = L'^\top D' R'$ be the singular value decompositions of H and H' . One has $H^+ = R^\top D^{-1} L$ and $H'^+ = R'^\top D'^{-1} L'$. One has:

Lemma 12. *Let H and $H' = H + E$ be two $n \times m$ matrices. Let $H = L^\top D R$ and $H' = L'^\top D' R'$ be the singular value decompositions of H and H' . Let σ be such that $\forall i, \sigma_i \geq \sigma, \sigma'_i \geq \sigma$. One has*

$$\|D^{-1} - D'^{-1}\|_F \leq \|D^{-1} - D'^{-1}\|_* \leq \frac{d\|E\|_2}{\sigma^2}$$

Proof. On has $|\frac{1}{\sigma_i} - \frac{1}{\sigma'_i}| \leq |\frac{\sigma'_i - \sigma_i}{\sigma_i \sigma'_i}| \leq \frac{\|E\|_2}{\sigma^2}$, and one has the conclusion. \square

The following result is straightforward from [19]:

Lemma 13. *Let H and $H' = H + E$ be two matrices. Let $\sigma_1 \geq \dots \geq \sigma_n$ be the singular values of H , and let $\sigma'_1 \geq \dots \geq \sigma'_n$ be the singular values of H' . Let σ be such that $\forall i, \sigma_i \geq \sigma, \sigma'_i \geq \sigma$. Let $H = L^\top D R$ and $H' = L'^\top D' R'$ be the singular value decompositions of H and H' . One supposes that $\|E\|_F \leq \sigma/2$. One then has*

$$\|L - L'\|_F \leq \frac{4(2\sqrt{d}\|H\|_F\|E\|_F + \|E\|_F^2)}{\sigma^2}, \|R - R'\|_F \leq \frac{4(2\sqrt{d}\|H\|_F\|E\|_F + \|E\|_F^2)}{\sigma^2}$$

7.4.1 Proof of Proposition 6

Let μ_ϵ be as in Lemma 8. The condition on N implies $\mu < \mu_\epsilon$. Let $H \in \mathcal{H}_\mu^S$, there exists $H' \in \mathcal{H}_0$ such that $\|H - H'\|_F < \epsilon$. One has $\|L\|_F = \|L'\|_F = \|R\|_F = \|R'\|_F = \sqrt{d}$, as the matrices are orthonormal. One has also $\|D^{-1}\|_F \leq \sqrt{d}/\sigma$. One uses the equality $AB - A'B' = (A - A')B - (A - A')(B - B') + A(B - B')$. One has

$$\begin{aligned} H^+ - H'^+ &= L^\top D^{-1} R - L'^\top D'^{-1} R' \\ &= L^\top [(D^{-1} - D'^{-1})R - (D^{-1} - D'^{-1})(R - R') + D^{-1}(R - R')] \\ &\quad - (L^\top - L'^\top)[(D^{-1} - D'^{-1})R - (D^{-1} - D'^{-1})(R - R') + D^{-1}(R - R')] + (L^\top - L'^\top)D^{-1}R \end{aligned}$$

Using the previous inequalities, and keeping only the first order terms, leads to

$$\|H^+ - H'^+\|_F \leq O\left(\frac{d^2 \epsilon}{\sigma^3}\right)$$

One also has $\|H^+\|_F \leq \frac{d^2}{\sigma}$. Plugging all those inequalities in the formulas computing the FSTs parameters leads to the result. \square